



PERSONAL DETAILS

EMAIL
saibalajinamburi@gmail.com

LOCATION
Hildesheim, Germany

LINKEDIN
linkedin.com/in/saibalajinamburi

GITHUB
github.com/saibalajinamburi

PORTFOLIO
saibalaji-portfolio.vercel.app

DATE OF BIRTH
17 February 2004

NATIONALITY
Indian

LANGUAGES

English IELTS 7 — Fluent

German A1 — Beginner
(actively learning, daily practice)

Telugu Native

TECH STACK

ML & DEEP LEARNING

Python PyTorch scikit-learn

LightGBM HuggingFace Pandas

MLOPS & DEPLOYMENT

MLflow DagsHub Docker

Kubernetes FastAPI CI/CD

GitHub Actions

RAG & AGENTIC AI

LangChain LangGraph LlamaIndex

ChromaDB FAISS

INFRA & MONITORING

PySpark Airflow Prometheus

Grafana Streamlit Web Scraping

SQL Git

METHODOLOGIES

Agile / Scrum System Design

A/B Testing Data Privacy

Problem Solving

Saibalaji Namburi

ML Engineer | MLOps | LLM Systems (RAG, Agentic AI)

PROFILE

ML Engineer with a background in building production-grade ML systems end-to-end — from model development and experiment tracking to containerised deployment, CI/CD pipelines, and live monitoring. Highly interested in exploring and implementing RAG architectures and agentic AI workflows, dedicated to bridging the gap between theoretical models and robust, scalable software.

CORE COMPETENCIES

- ML Model Development & Training
- MLOps Pipeline Design & CI/CD for ML
- Model Deployment & API Engineering
- Production Logging & Monitoring
- RAG Systems & Vector Databases
- Agentic AI Architectures & LLM Systems

EDUCATION

- Apr. 2026 – Present **M.Sc. International Data Analytics**
University of Hildesheim — Germany
Focus: Machine Learning, MLOps, AI Systems, Data Pipelines
- Aug. 2021 – Jul. 2025 **B.Tech Artificial Intelligence & Machine Learning**
Saveetha School of Engineering — India
GPA: 8.756 / 10 · Focus: Deep Learning, NLP, Data Processing

EXPERIENCE

- 04.2024 — 03.2025 **AI Application Developer — Internship**
Saveetha School of Engineering, Chennai
- Designed and launched DentalSolutions on Google Play — a patient-doctor platform with appointment booking and WhatsApp integration
 - Owned backend API development (FastAPI/Python) and implemented automated monitoring workflows for production-grade availability
 - Optimised database queries, improving data retrieval speed by over 35%
 - Delivered ML feature deployments with quality assurance and A/B testing

PROJECTS

- MLOPS · GENAI **SupportPulse — Enterprise AI Support Triage Platform**
- Architected a multi-layer AI system: LLM Cascade for classification, LightGBM for SLA prediction, and ChromaDB for semantic retrieval
 - Implemented a local-first LLM strategy using Gemma2:2b and Gemma4:e4b fallback to ensure data privacy and high accuracy
 - Built a production-grade MLOps lifecycle: MLflow tracking, Feast Feature Store, and Great Expectations data validation
 - Established automated CI/CD (GitHub Actions) for Docker deployments, and real-time observability (Prometheus/Grafana)
 - Processed 68,000+ tickets with 2.08ms ChromaDB retrieval latency, achieving 1.0 RAGAS Faithfulness and 100% LLM Cascade accuracy, validated by 38 Pytest unit tests (1.87s runtime)
- MLOPS · NLP **ResearchIQ — Automated Research Insight Engine**
- Processed 42k+ scientific abstracts, classifying them into 15 domains using targeted Regex NLP cleaning
 - Trained a Scaled Logistic Regression (One-vs-Rest) model for multi-label abstract classification
 - Engineered an ONNX inference engine (INT8 MiniLM), cutting memory by 80% and reducing latency to <10ms
 - Integrated DVC for large dataset versioning and DagsHub/MLflow for remote experiment tracking
 - Deployed containerized FastAPI and Streamlit surfaces to Hugging Face Spaces
 - Achieved a competitive 0.6655 Global Macro F1 score across 15 overlapping academic domains, maintaining 100% label parity post-ONNX conversion